

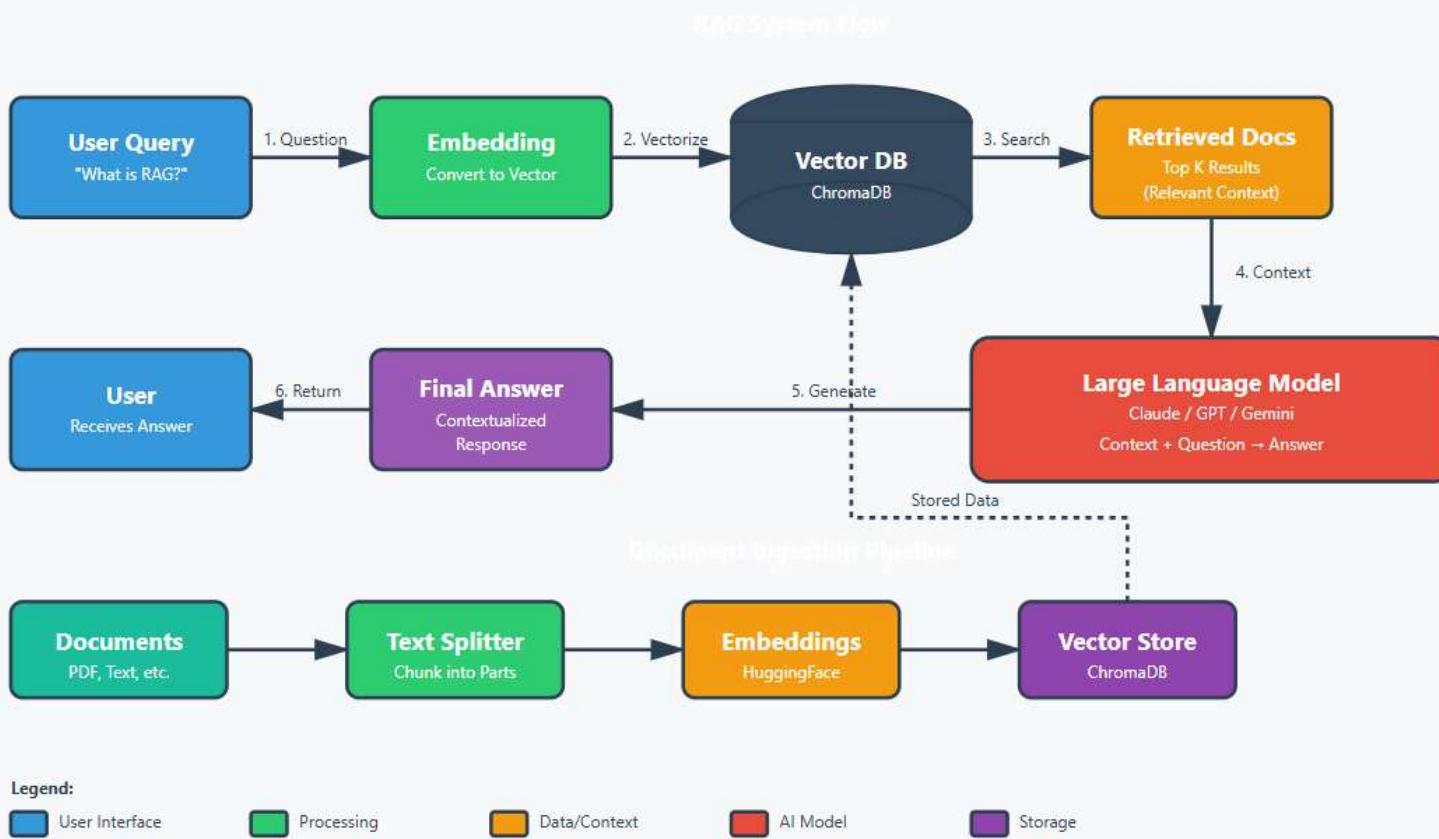
# RAG MCP Server

With ChromaDB, LangChain, Claude LLM

Lucas Kim

November 26, 2025

# RAG (Retrieval-Augmented Generation) Architecture





## RAG Process Flow

### 1 Document Ingestion (add\_documents)

Receive source documents → Split into chunks using RecursiveCharacterTextSplitter → Convert to vectors using HuggingFace embeddings → Store in ChromaDB with metadata

### 2 Vector Search (search\_documents)

Convert user query to embedding → Search ChromaDB using cosine similarity → Return top k documents (default: 4)

### 3 Context Construction

Merge retrieved documents into unified context → Maintain document order and source information → Generate prompt template

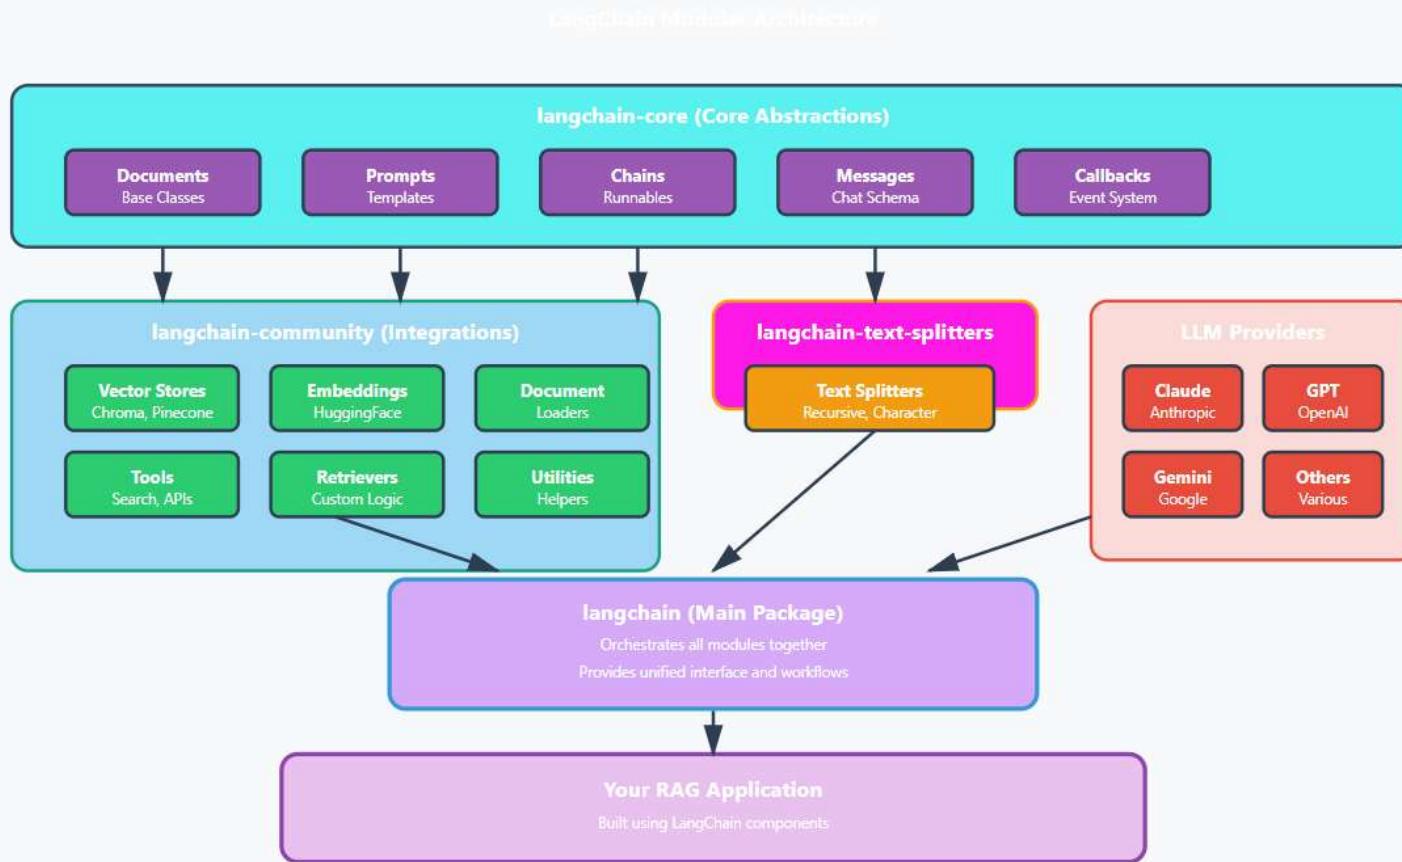
### 4 AI Answer Generation (rag\_query)

Send context + question to Claude API → Claude Sonnet 4 model generates answer → Return result with source citations

### 5 Result Formatting & Return

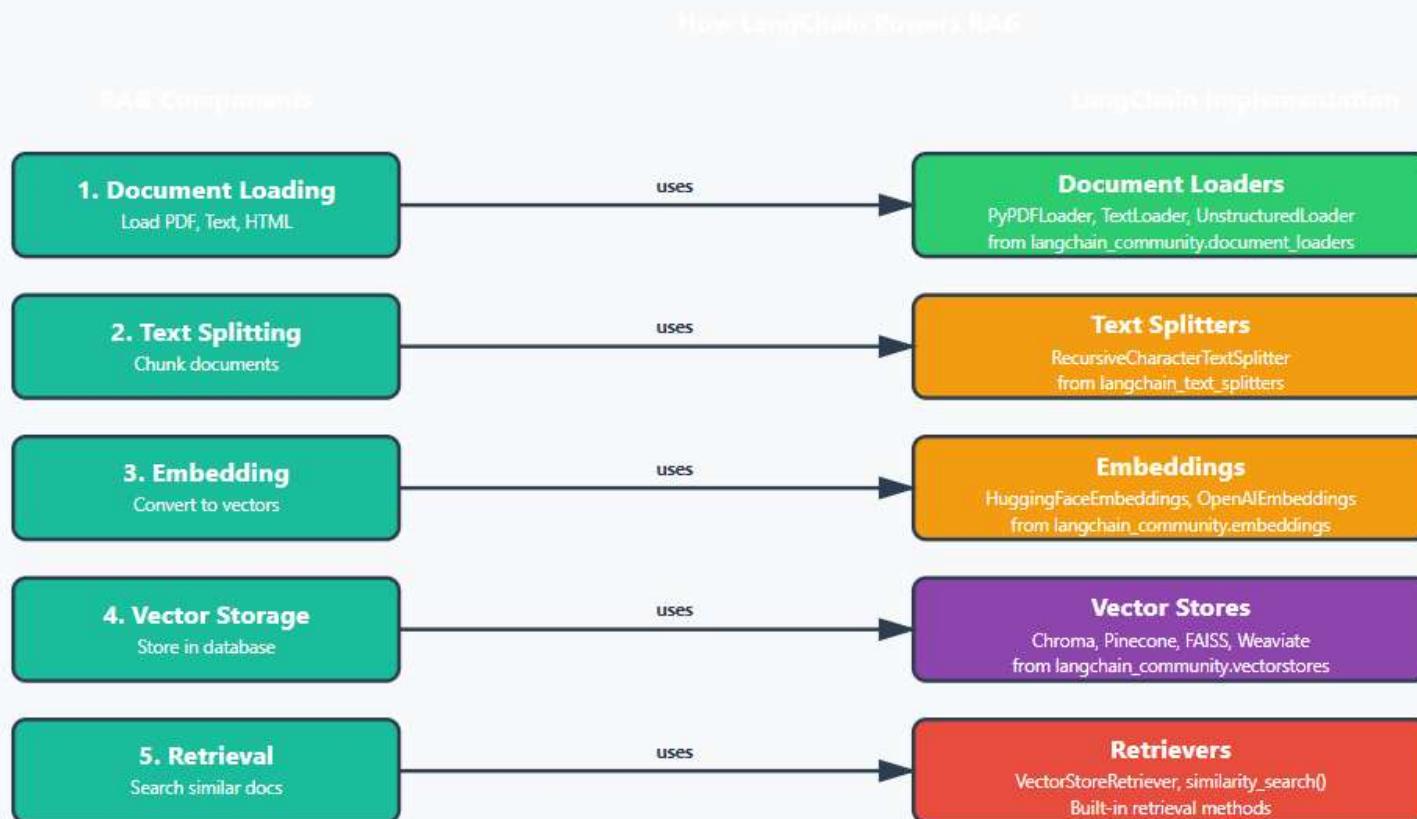
Format answer text + reference document info → Include metadata (source, chunk position, etc.) → Deliver final result to user

## LangChain 1.0 Architecture



Module Dependencies: ➡ Depends on   ■ Main Package   ■ Community Integrations   ■ LLM Providers   ■ Core Abstractions

## 🔗 RAG + LangChain Integration



## System Architecture



# Download source

```
#Download source
```

```
$> git clone https://github.com/hjkim7796/RAG-CHROMA-MCP-SERVER.git
```

```
$> dir RAG-CHROMA-MCP-SERVER
```

```
Directory: C:\hjkim\PythonProject\RAG-CHROMA-MCP-SERVER
```

Mode	LastWriteTime	Length	Name
d----	11/21/2025 12:06 PM		.venv
d----	11/21/2025 11:58 AM		docs
-a---	11/21/2025 11:58 AM	68	.gitignore
-a---	11/21/2025 11:58 AM	227607	abc.pdf
-a---	11/21/2025 11:58 AM	16847	add_pdf_to_mcp.py
-a---	11/21/2025 11:58 AM	5657	mcp-http-proxy.py
-a---	11/21/2025 11:58 AM	383	query-chromaDB.py
-a---	11/21/2025 11:58 AM	20699	rag_mcp_http_server.py
-a---	11/21/2025 11:58 AM	7821	README_HTTP.md
-a---	11/21/2025 11:58 AM	347	requirements.txt
-a---	11/21/2025 11:58 AM	6936	requirements_freeze.txt
-a---	11/21/2025 11:58 AM	11420	test_server.py
-a---	11/21/2025 11:58 AM	2034	test_server_with_proxy.py

# Run rag-mcp-server

```
#Go to the source folder
```

```
$> cd <user-project-folder>
```

```
#Config a virtual environment (venv) in VS Code for Python to ensure your project uses the correct isolated environment.
```

```
$> python -m venv .venv
```

```
$> .venv/Scripts/activate.ps1
```

```
#Install all python packages
```

```
(.venv)> pip install -r requirements.txt
```

```
#Run rag mcp server
```

```
(.venv)> python rag_mcp_http_server.py
```

# Test rag-mcp-server

```
#Go to the source folder
```

```
(.venv)> cd <user-project-folder>
```

```
#Check server is working properly
```

```
(.venv)> python test_server.py
```

```
#Check proxy server is working properly
```

```
(.venv)> python test_server_with_proxy.py
```

# Config for Claude Desktop integration

The screenshot shows the Claude Desktop application's settings interface. On the left, a sidebar lists various settings categories: General, Account, Privacy, Billing, Usage, Capabilities, Connectors, Claude Code, Desktop app, General, Extensions, and Developer. The 'Developer' category is currently selected, indicated by a highlighted background. To the right, a main panel titled 'Local MCP servers' displays a list of configured servers. One server, 'rag-search', is highlighted with a light gray background. An 'Edit Config' button is located at the top left of this panel. Below the list, detailed configuration information is shown for the 'rag-search' server, including its command path (C:\hjkim\PythonProject\RAG-CHROMA-MCP-SERVER\venv\Scripts\python) and arguments (C:\hjkim\PythonProject\RAG-CHROMA-MCP-SERVER\mcp-http-proxy.py http://localhost:8000/sse). A blue arrow points from the 'rag-search' entry in the list to the corresponding JSON configuration code below.

```
File location: %APPDATA%\Claude\claude_desktop_config.json
{
  "mcpServers": {
    "rag-search": {
      "command": "C:\hjkim\PythonProject\RAG-CHROMA-MCP-SERVER\venv\Scripts\python",
      "args": ["C:\hjkim\PythonProject\RAG-CHROMA-MCP-SERVER\rag_mcp_server.py"]
    }
  }
}
```

# Test rag-mcp-server

```
#Go to the source folder
```

```
(.venv)> cd <user-project-folder>
```

```
#Check server is working properly
```

```
(.venv)> python test_server.py
```

```
#Check proxy server is working properly
```

```
(.venv)> python test_server_with_proxy.py
```

# Add a pdf document

```
#Go to the source folder
```

```
(.venv)> cd <user-project-folder>
```

```
#Add pdf document to RAG system
```

```
(.venv)> python add_pdf_to_mcp.py abc.pdf
```

```
...
```

```
=====
```

PDF successfully added to RAG system!

```
=====
```

# Test query on claude desktop

Query: **Connect to the rag-search MCP server and summarize "EO Patch Wearable Insulin Pump"**

L Connect to the rag-search MCP server and summarize "EO Patch Wearable Insulin Pump"

I'll connect to the rag-search MCP server and search for information about the "EO Patch Wearable Insulin Pump".

R Rag query

It looks like there's a configuration issue - the ANTHROPIC\_API\_KEY isn't set up for the RAG query function. Let me try searching the documents directly instead:

R Search documents

Great! I found information about the EO Patch Wearable Insulin Pump. Let me get more details:

R Search documents

Perfect! Based on the documents in the RAG database, here's a summary of the **EO Patch Wearable Insulin Pump**:

### Overview

The EO Patch is an external, wearable insulin injection device (also called a "patch pump") that automatically delivers insulin from outside the body to control blood glucose levels.

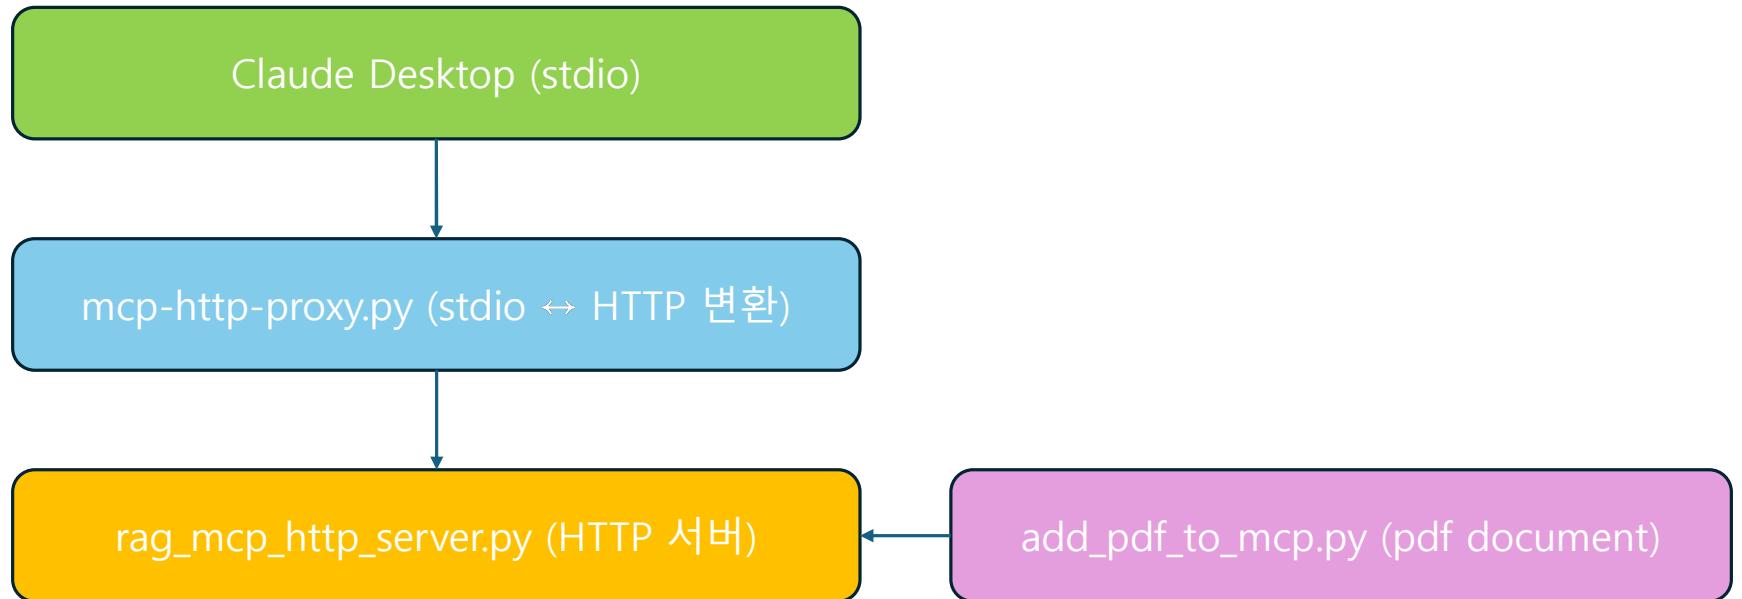
To enable Rag queries, create a .env file and add the following:  
ANTHROPIC\_API\_KEY=your-api-key-here



# File Description

- [rag\\_mcp\\_server.py](#) - Main RAG MCP server
- [add\\_pdf\\_to\\_mcp.py](#) – Example for adding a PDF document to RAG system
- [query-chromaDB.py](#) – Example for querying to chromaDB
- [test\\_server.py](#) - Example for testing the RAG MCP server
- [test\\_server\\_with\\_proxy.py](#) - Example for testing the Proxy server
- [requirements.txt](#) – All python package list
- [README\\_HTTP.md](#) - Detailed User Guide

# How it works





# Key Features

- **add\_documents** - Adding documents to the vector DB
- **search\_documents** - Similarity-based search
- **rag\_query** - Query via RAG (using Claude), Need API key
- **delete\_collection** - Delete collection

# Available MCP Tools

## 1 add\_documents

Add documents to the vector database.  
Automatically handles chunk splitting and embedding.

**Parameters:**

- texts: Array of document texts
- metadatas: Array of metadata objects
- chunk\_size: Chunk size (default: 1000)
- chunk\_overlap: Overlap size (default: 200)

## 2 search\_documents

Search for similar documents in the vector database using cosine similarity.

**Parameters:**

- query: Search query string
- k: Number of results (default: 4)
- filter: Optional metadata filter

## 3 rag\_query

Answer questions using RAG. Retrieves relevant documents and generates answers with Claude.

**Parameters:**

- question: Question text
- k: Number of docs to search (default: 4)
- language: Answer language (ko/en)

## 4 get\_collection\_info

Retrieve current collection information including document count and storage location.

**Parameters:**

- None (no parameters required)

## 5 delete\_collection

Delete entire collection. All documents and embeddings will be permanently removed.

**Parameters:**

- confirm: Confirmation flag (must be true)

## Key Features

### Automatic Text Chunking

Documents are automatically split into manageable chunks (1000 chars with 200 char overlap) while preserving semantic coherence.

### Similarity Search

Cosine similarity-based vector search efficiently finds the most relevant documents for any query.

### Smart Embeddings

Uses sentence-transformers/all-MiniLM-L6-v2 model for high-quality embeddings. Runs locally without API keys.

### Metadata Management

Store and filter documents by custom metadata (source, category, tags, etc.) for precise retrieval.

### Persistent Storage

ChromaDB stores vectors persistently on disk. Data survives server restarts and can be queried anytime.

### Multilingual Support

Supports multiple languages through HuggingFace embeddings and Claude's multilingual capabilities.



## Technology Stack

---

**LangChain**

v1.0.7

**ChromaDB**

v1.3.5

**Claude AI**

Sonnet 4

**HuggingFace**

Transformers v5.1.2

**MCP**

v1.21.2

**Python**

v3.12